

AI-Augmented Project Controls for Large Capital Projects

Prof. Dr. Johannes Meier

June 2026

Date: June 2026

Executive Summary

Large capital projects systematically underperform. Across a database of 16,000 projects, Oxford University researchers found that only 8.5 percent met their original cost and schedule targets — and a mere 0.5 percent delivered all promised benefits on time and on budget.¹ McKinsey’s analysis of more than 300 projects exceeding one billion dollars in value found average cost overruns of 80 percent and schedule delays of 50 percent.² These are not outlier statistics; they describe the norm.

The root cause is not a shortage of engineering talent or project management methodology. It is an information problem: project controls data is generated too slowly, integrated too rarely, and interpreted too inconsistently for management to act before damage accumulates.

This paper proposes a practical, deployable response: a structured AI scaffolding system built on top of Claude Code — Anthropic’s enterprise AI agent — that converts raw project data into structured cost performance analysis, escalation-ready reports, and visual dashboards in minutes rather than days. The system is not a replacement for experienced project professionals. It is a force multiplier: it ensures that earned value calculations are performed consistently every reporting period, that escalation thresholds are enforced automatically, and that domain heuristics from decades of benchmarking are applied without relying on institutional memory.

The system has been designed, implemented, and tested on a hypothetical 500 km hydrogen transmission pipeline (EUR 1.75 billion CAPEX). It is available as open-source code at github.com/jmeier1963/large_capital_project_management.

This paper asks the board to authorize a pilot on one live capital project, to assign a senior executive as accountable owner, and to commission a data governance policy that defines which

¹Flyvbjerg, B. & Gardner, D. (2023). *How Big Things Get Done*. Macmillan. Database of 16,000 projects; 8.5% delivered on cost and schedule; 0.5% delivered all promised benefits.

²McKinsey & Company (2015). “Megaprojects: The good, the bad, and the better.” Review of 300+ projects exceeding USD 1 billion.

project data is permissible to process through AI systems and under what conditions.

1. The Capital Project Performance Crisis and the AI Opportunity

1.1 The Scale of the Problem

Flyvbjerg's landmark study — the largest empirical analysis of project performance ever conducted — established what he calls the “iron law of megaprojects”: over budget, over time, over and over again.³ The headline numbers bear repeating:

- **91.5%** of projects exceed their original budget or schedule, or both
- **8.5%** deliver within original parameters
- **0.5%** deliver cost, schedule, *and* promised benefits

McKinsey adds sectoral granularity. Rail projects overrun by an average of 45 percent. Bridges and tunnels by 35 percent. Mining and metals projects — among the most capital-intensive category — see 83 percent of major projects exceed planned CAPEX by more than 40 percent, with average delays of 20 to 30 months.⁴

The energy transition amplifies the stakes. The Hydrogen Council and McKinsey reported in 2025 that the global hydrogen sector has now committed over USD 110 billion across more than 500 projects past Final Investment Decision.⁵ These are projects where schedule delays translate directly into stranded capital, missed carbon commitments, and competitive disadvantage in a market where first-mover timing is commercially decisive.

1.2 Why Projects Fail: The Information Gap

The intuition that projects fail because of poor engineering or inadequate planning is only partially correct. McKinsey's analysis of 48 deeply troubled megaprojects found that **73 percent of cost and schedule overruns were caused by poor execution** — specifically, by failures of monitoring, escalation, and corrective action — not by flawed original designs.⁶

The mechanism is predictable. Project data — cost actuals, schedule progress, change orders, risk events — is generated continuously in the field. But in most organisations, this data travels slowly: from site time-sheets and contractor invoices into cost control systems, then into spreadsheets maintained by cost engineers, then into narrative reports assembled for manage-

³Flyvbjerg, B. (2017). “The Iron Law of Megaprojects.” Cato Institute Policy Report. 91.5% of projects exceed budget, schedule, or both.

⁴McKinsey Global Institute (2016). “Reinventing Construction.” Mining and metals sector overrun data. Rail, bridge, and tunnel statistics from Flyvbjerg's database.

⁵Hydrogen Council & McKinsey & Company (2025). *Global Hydrogen Compass*. USD 110B committed investment across 500+ post-FID projects.

⁶McKinsey & Company (2017). “Increasing transparency in megaproject execution.” Analysis of 48 troubled projects; 73% of overruns attributed to execution failures.

ment review. By the time a negative trend is visible in a board report, the underlying pattern may have been developing for three or four months.

The consequence is that corrective actions are taken late, when the cost of recovery is high. A schedule slip that costs EUR 2 million to correct in month three may cost EUR 20 million by month nine — not because the underlying problem grew tenfold, but because the window for low-cost intervention closed.

1.3 The Project Controls Baseline Problem

Consistent, rigorous project controls — earned value management (EVM), integrated schedule analysis, systematic risk quantification — are the established antidote. The Independent Project Analysis Group (IPA), whose proprietary database covers more than 20,000 capital projects, measures project performance against a Project Control Index (PCI). Projects with strong project controls consistently outperform their peers on cost, schedule, and operability outcomes.⁷

The barrier to universal adoption of rigorous project controls is not lack of awareness. It is capacity and consistency. A cost engineer maintaining full earned value analysis across eight WBS packages, three active contracts, and five reporting periods simultaneously — while also managing contractor queries, validating invoices, and preparing the monthly report — will, under pressure, abbreviate the analysis. Thresholds will not be checked. EAC forecasts will not be run across all three methods. The escalation that should have gone to the Project Director this week will go next week instead.

This is the specific gap that AI-augmented project controls addresses.

1.4 What Has Changed: The AI Inflection Point

For most of the past decade, AI in project management meant dashboards that required extensive data pipeline engineering, natural language processing tools that extracted text with limited accuracy, and predictive models that required months of proprietary training data before producing actionable output. The technology was real but the deployment cost was prohibitive for all but the largest programmes.

The shift since 2023 is qualitative, not incremental. Large language models (LLMs) — in particular, AI agents that reason over structured data, apply domain-specific rules, call tools, and generate structured outputs — have changed what is achievable without bespoke development.

Three capabilities now combine in a way that was not previously available:

Structured reasoning over domain rules. An AI agent can be given a plain-English rules file (e.g., “if CPI drops below 0.85, generate a Project Director escalation memo; if SPI has been below 0.85 for two consecutive periods, issue a mandatory written notice”) and will apply those

⁷Independent Project Analysis (IPA). *Capital Project System Improvement*. Project Control Index methodology and performance correlation. ipaglobal.com

rules with perfect consistency on every run. The rules are not remembered from a training corpus — they are encoded explicitly and executed deterministically.

Tool use and computation. Modern AI agents can invoke Python scripts, read CSV files, write structured output, and generate charts as a native part of their workflow. The EVM module in this system runs a Python engine (ANSI/EIA-748 compliant, 713 lines) that computes ten financial metrics per WBS package, runs three EAC forecast methods, produces five publication-ready charts, and writes a complete markdown report — all as a single AI-orchestrated pipeline.

Domain heuristics without domain re-training. Parametric cost benchmarks, productivity norms, and contingency ranges can be loaded into the AI's context at runtime as versioned YAML files, without fine-tuning or model retraining. When the AI applies a benchmark, it cites the source and version. When project data falls outside the benchmark's valid range, it flags this explicitly. The heuristics library is updated as the organisation accumulates actuals — turning project experience into institutional knowledge that the AI accesses on the next run.

This is not a claim that AI understands project management in the way an experienced cost engineer does. It is a more precise and more practically important claim: AI can perform the *routine, structured, rule-governed portion* of project controls work with perfect consistency, freeing skilled professionals to focus on the *judgment-intensive* portions — contractor negotiation, root cause analysis, recovery planning — where human expertise is irreplaceable.

Avoiding the enterprise AI trap. Deloitte's 2025–2026 State of AI survey found that while nearly 90 percent of companies have deployed AI in at least one business function, 94 percent report not seeing significant value from their investments.⁸ The explanation is structural: most AI deployments target individual productivity (drafting emails, summarising documents) rather than the *process* level where value is concentrated.

In capital project controls, value is concentrated in process consistency, not individual productivity. The question is not whether a cost engineer can produce a report faster with AI assistance. The question is whether every project — regardless of the experience level of its controls team, regardless of whether it is month four or month thirty-two — produces a complete, consistently structured, threshold-checked EVM report on time, every reporting period.

The approach described in this paper embeds AI into the process as a non-optional execution step, with encoded domain rules, parameterised thresholds, and mandatory output formats that cannot be abbreviated under pressure.

2. The System: Architecture and Capabilities

⁸Deloitte US (2026). *State of AI in the Enterprise*. 89% of companies have deployed AI; 94% report not seeing significant value.

2.1 Design Principles

The system is built on four principles that distinguish it from generic AI implementations:

Data first, conversation second. Project intelligence is derived from structured data — CSV files, YAML configuration, JSON schemas — not from narratives or slide decks. This makes outputs reproducible, auditable, and comparable across periods.

Provenance on every heuristic. Cost benchmarks, productivity norms, and contingency ranges are stored as versioned YAML files with mandatory `source:` and `valid_range:` fields. When the AI applies a benchmark, it cites the source. When the project falls outside the valid range, it flags this explicitly.

Rules encoded, not remembered. Escalation thresholds, required approvals, and judgment rules are encoded in a `CLAUDE.md` configuration file — the project’s AI constitution. They are applied automatically on every analysis run, not recalled from a briefing document.

Roles, not generics. Every communication or action recommendation produced by the AI addresses a specific named role from the project’s organisation chart. It does not recommend that “the team” should act. It specifies that the Commercial Manager should review the change order, that the Project Director should receive the escalation memo, and that the cost register should be updated by the Cost Engineer within five business days.

2.2 System Components

The system consists of five integrated layers.

Layer 1 — Project Data Store

A structured directory of CSV, YAML, JSON, and Markdown files representing the project’s living state. The recommended directory layout mirrors standard project controls practice:

```
my-project/
  CLAUDE.md           ← project AI constitution
  project-brief.md    ← scope, FID CAPEX, key milestones
  01-contracts/
    contracts-register.csv
    change-orders/
  02-schedule/
    milestones.csv     ← milestone_id, planned_date, forecast_date, critical_path
  03-cost/
    evm-timephased.csv ← one row per WBS element per reporting period
    evm-output/        ← written by the EVM skill each period
  04-resources/
    resource-matrix.csv
```

05-risk/

risk-register.csv

All files are validated against JSON schemas on intake. The directory is version-controlled, meaning every change is logged and reversible.

The CLAUDE.md AI Constitution. The central configuration file that defines how the AI must behave on this project. A production CLAUDE.md for a capital project encodes mandatory judgment rules in plain English:

- *“If CPI drops below 0.85, immediately generate a Project Director escalation memo referencing the WBS package, the current CPI value, and the three-period CPI trend.”*
- *“If SPI has been below 0.85 for two consecutive reporting periods, generate a mandatory written escalation notice. Do not wait for a third period.”*
- *“Any change order exceeding 1% of the contract value triggers an independent cost review. Notify the Commercial Manager and the PMO lead.”*
- *“At stage-gate FEED-to-EPC, the estimate must be AACE Class 3 or better. Flag any estimates submitted at Class 4 or 5 as non-compliant.”*

It also defines stage-gate-specific standards, so the same system applies different rigor depending on the project phase:

Phase	Estimate Class	Contingency Range	Schedule Basis
Concept	Class 5	35–40%	±50%
Pre-FEED	Class 4	25–30%	±30%
FEED	Class 3	15–20%	±15%
Execution	Class 2	8–12%	±10%
Closeout	Class 1	3–5%	Actuals

These rules are applied on every AI run without exception. They cannot be abbreviated because the cost engineer is under pressure.

Layer 2 — Domain Heuristics Library

A curated library of parametric benchmarks and productivity norms, stored as versioned YAML files. Two files are included for hydrogen pipeline projects:

heuristics/pipeline-cost.yaml — installed cost per kilometre by pipe diameter and terrain type (abbreviated):

```
# Source: IPA Benchmarking Database + DVGW project actuals (Western Europe)
# Valid for: onshore H2 pipelines, Germany / Western Europe, AACE Class 3–5
# Accuracy: ±20–30% (Class 3 equivalent)
# Version: 1.0 | Updated: 2026-01
```

```
base_cost_eur_per_km:
  DN300: 1_450_000
  DN400: 1_920_000      # H2-PIPE-DE-001 reference value
  DN500: 2_650_000
  DN600: 3_400_000

terrain_factors:
  flat_agricultural: 1.00
  rolling_rural:      1.15
  urban_corridor:     1.60
  river_crossing_hdd: 2.50  # per crossing

h2_service_premium: 0.15  # +15% for H2-grade materials (HIC-tested, dry-gas seals)

contingency_by_class:
  class_5: [0.30, 0.50]  # conceptual estimate
  class_3: [0.15, 0.25]  # study estimate
  class_2: [0.10, 0.15]  # budget-quality estimate
  class_1: [0.05, 0.10]  # definitive estimate
```

When the AI applies this benchmark to a cost estimate, it cites the source and version, and flags any project falling outside the valid range (e.g., a DN900 pipeline or a project in a different region). `heuristics/productivity-norms.yaml` covers labour productivity for pipeline welding (joints per shift by diameter and welder grade), civil excavation (m³/shift by terrain), mechanical erection, and E&I installation, with the same provenance structure.

Layer 3 — Specialist Agents

Four domain-specific agents are defined as markdown files in `scaffolding/agents/`. Each agent specifies its trigger conditions, analysis steps, output format requirements, and escalation rules. They are installed by copying to `.claude/agents/` in the project directory, where Claude Code discovers and invokes them automatically on matching prompts.

Agent	Trigger	Output
EVM Analyst	Monthly cycle; cost or schedule performance data	EV report, 5 charts, 3 EAC forecasts, RAG status, escalation flags
Schedule Analyzer	Milestone slip; SPI below threshold; “schedule review”	Schedule health report, critical path float, recovery options

Agent	Trigger	Output
Contract Reviewer	Change order; claim notice; new contract	Commercial summary, flagged clauses, CO entitlement assessment
Risk Assessor	Monthly cycle; new risk identified	EMV-ranked Top-10 risk digest, escalation flags

The `schedule-analyzer` agent executes a defined sequence: read `milestones.csv` and flag any milestone with a forecast slip exceeding 14 days; calculate total float on the critical path from the latest schedule export; compare the current SPI trend against the project's historical SPI curve; generate a recovery options matrix with cost and schedule impact for each option. Its escalation rule is encoded in its definition: any slip greater than 60 days on a milestone marked `critical: true` triggers a mandatory Project Director memo, regardless of whether the SPI threshold has been breached.

The `contract-reviewer` agent processes change orders against a structured extraction template: contract type, base value, CO value as a percentage of contract, entitlement basis (scope change vs. changed conditions), dispute resolution mechanism, and liquidated-damages exposure. It cross-references the CO value against the `CLAUDE.md` threshold and generates the independent cost review notification automatically when the threshold is exceeded.

Layer 4 — Claude Code Skills

Claude Code's skill system packages purpose-built Python tools for automatic invocation. Skills are installed by placing a directory in `~/.claude/skills/`, containing a `SKILL.md` behavioural instruction file and the Python implementation. Claude Code scans this directory at startup and invokes skills when user prompts match the trigger phrases defined in `SKILL.md`.

Skill	Source	Capability
<code>evm</code>	Purpose-built	EVM engine: 10 metrics, 5 charts, 3 EAC forecast methods
<code>budget-estimator</code>	Purpose-built	Parametric CAPEX with P50/P90 Monte Carlo simulation
<code>pdf</code>	Marketplace	Extract structured text from contracts and specifications
<code>pptx</code>	Marketplace	Assemble progress presentation from EVM report and charts
<code>csv-data-summarizer</code>	Marketplace	Statistical summary of cost registers and change order logs

Skill	Source	Capability
meeting-insights-analysis	Marketplace	Extract action items and decisions from meeting notes

The EVM skill triggers on “EVM”, “earned value”, “CPI/SPI”, or any CSV with columns bac, bcws_cum, bcwp_cum, acwp_cum. It auto-detects snapshot vs. time-phased format and generates a complete ANSI/EIA-748 report in under two minutes.

The budget-estimator skill triggers on “estimate CAPEX”, “P50/P90”, “budget estimate”, or “what will this cost”. It accepts a YAML work-package definition or a set of CLI parameters and runs 10,000 Monte Carlo iterations using triangular distributions calibrated to the AACE estimate class selected. Each work package can be specified parametrically (using the pipeline-cost heuristics: EUR/km by diameter and terrain, H2 service premium, compression station cost per MW) or as a direct cost entry with an uncertainty range.

The output is an AACE-classified estimate report with three values: P10 (optimistic boundary), P50 (central planning estimate), and P90 (risk-adjusted ceiling for budget approval). A sensitivity tornado chart ranks the work packages by their contribution to the P90–P10 spread, identifying where contingency is being driven and where engineering definition would most reduce uncertainty. On the hydrogen pipeline example, the Class 3 estimate produces a P50 of EUR 1,833M and a P90 of EUR 1,930M against the approved FID CAPEX of EUR 1,749M — consistent with the $\pm 20\%$ accuracy band of a FEED-stage estimate.

Layer 5 — JSON Schemas and Validation

Three JSON Schema files enforce data quality at intake:

- `schemas/project-brief.schema.json` — validates the YAML frontmatter of `project-brief.md`: required fields (`project_id`, `fid_date`, `approved_capex_eur`, `contingency_eur`), data types, and allowable values for `status` and `phase`.
- `schemas/contract.schema.json` — validates contract register entries: contract type (lump sum / EPCM / reimbursable / time-and-materials), required date fields, and mandatory completion of `ld_rate_per_day` for lump-sum contracts.
- `schemas/risk-register.schema.json` — validates risk register entries: probability (0–1 float), impact (EUR value), required mitigation owner, and mandatory EMV recalculation trigger when either probability or impact changes by more than 10 percent.

Schema validation runs automatically when Claude Code opens a project directory. Any invalid data file generates an immediate quality flag before analysis begins.

2.3 The EVM Module: From Calculation to Institutional Discipline

What Earned Value Management Is — and Why It Is Underused

Earned Value Management is the internationally recognised standard (ANSI/EIA-748) for integrating cost and schedule performance measurement. Its core logic is simple: at any point in a project, you can compute not just what you have spent (Actual Cost, ACWP), but what you *should* have spent for the work you have actually completed (Earned Value, BCWP). The ratio of earned value to actual cost is the Cost Performance Index (CPI). A CPI of 0.91 means you are spending EUR 1.10 to deliver EUR 1.00 of budgeted work.

The power of EVM lies in its predictive validity. Research across thousands of projects has established that the CPI at the 20 percent completion milestone is a reliable predictor of final outcome. Projects that are underperforming at 20 percent completion rarely recover to budget — and when they do, it is because management intervened early, not because efficiency spontaneously improved.⁹

Despite this, EVM is performed inconsistently in practice. The calculation requires integrating cost actuals, progress measurements, and the original budget baseline — data that typically live in three or four separate systems and require manual reconciliation. Under schedule pressure, this reconciliation gets abbreviated. The result is that the single most powerful early warning indicator available to project management is produced late, inconsistently, or not at all.

What the EVM Module Delivers

Given a CSV file containing WBS codes, budget at completion (BAC), planned value (BCWS), earned value (BCWP), and actual cost (ACWP) — data that any functioning cost control system produces — the EVM module generates, in under two minutes:

- **A complete markdown report** with executive summary, WBS-level performance table with RAG status, three EAC forecasts (CPI method, composite CPI×SPI method, and planned-rate method), variance at completion, and TCPI
- **Five publication-ready charts:** S-curve (BCWS/BCWP/ACWP over time), CPI/SPI trend with threshold bands, WBS cost variance waterfall, EAC forecast comparison, and a CPI/SPI quadrant map with bubble sizes proportional to budget at stake
- **Automatic escalation flags** when CPI or SPI breach thresholds, when SPI has been below 0.85 for two consecutive periods, or when TCPI exceeds 1.10 (the level at which recovery is generally considered unrealistic without re-baselining)
- **A plain-language interpretation** identifying the worst-performing package, the recommended EAC to use for board reporting, and the next three actions with named responsible parties

The RAG thresholds applied are industry-standard: $\text{CPI or SPI} \geq 0.95 = \text{GREEN}$; $0.85\text{--}0.94$

⁹Christensen, D.S. (1993). “The Estimate at Completion Problem: A Review of Three Studies.” *Project Management Journal*. CPI stability at 20% completion and its predictive validity for final outcome.

= AMBER; < 0.85 = RED. The composite CPI×SPI EAC method is recommended for board reporting, because it accounts for schedule pressure compounding cost efficiency loss — the dominant failure pattern in large capital projects.

For the hydrogen pipeline example included with the system, this method produces an EAC of EUR 1.637 billion against an approved BAC of EUR 1.350 billion for the contracted scope — a 21 percent overrun signal at month five of a 36-month execution programme, early enough for effective intervention. WBS-1.1 (Mainline North) is identified as the primary driver: its SPI of 0.750 (RED) reflects a linepipe delivery delay caused by port congestion, and if SPI remains below 0.85 in the June reporting period, the two-consecutive-period escalation rule fires automatically, generating a mandatory written notice to the Project Director.

2.4 Integration with Existing Tools

The system reads data that capital project organisations are already producing; it does not require replacing any existing tool.

Primavera P6. Schedule data is ingested via .xer export files using the P6XER MCP (Model Context Protocol) server. This allows Claude Code to read P6 project schedules directly, extract critical path float, identify milestone forecast dates, and feed schedule data into the schedule-analyzer agent — without requiring a P6 licence or database connection in the AI environment.

Excel and CSV cost systems. Any cost control system that can export EVM data in tabular form (BCWS, BCWP, ACWP by WBS code) is compatible. The EVM skill accepts both snapshot and time-phased formats and validates column names on intake.

Document repositories. Contract PDFs, engineering specifications, and tender documents are processed via the pdf skill. The contract-reviewer agent combines PDF extraction with structured templates to produce commercial summaries from documents that would otherwise require hours of manual review.

Reporting platforms. The pptx skill converts the monthly EVM markdown report and the five generated chart PNGs into a presentation-ready deck, formatted to the organisation's template. This output feeds directly into the monthly board reporting pack without manual transcription.

Atlassian and collaboration tools. The Atlassian MCP server enables action items generated by the AI (e.g., from the risk-assessor or schedule-analyzer agents) to be logged directly as Jira tickets with assigned owners and due dates, completing the loop from AI analysis to trackable task.

3. Investment Case, Risks, and Governance

3.1 The Value Drivers

The business case for AI-augmented project controls operates on three distinct value levers:

Lever 1 — Earlier detection of adverse trends

Research by IPA establishes that the cost of corrective action increases approximately exponentially with the delay between signal and response. A schedule slip that can be corrected for EUR 2 million in month three may require EUR 15–20 million to recover by month nine, because interim procurement commitments, subcontractor mobilisation costs, and lost float compound the original deviation.

The system makes the trend visible in the first reporting period it appears, rather than the third. On a EUR 1 billion project, even a single such early intervention — preventing a 2 percent cost growth that would otherwise materialise — generates a return that exceeds the entire cost of implementing and operating the system.

Lever 2 — Consistent application of escalation rules

Escalation failures are documented as a primary driver of late-stage project crises. The mechanism is well understood: a cost engineer identifies a negative trend, judges it borderline, decides to wait one more period to confirm before escalating, and the window for low-cost intervention closes. Encoded escalation thresholds eliminate the discretion at the margin. When CPI falls below 0.85, the flag fires automatically. The Project Director is notified regardless of where the cost engineer’s judgment falls.

Lever 3 — Reduced cost of producing controls outputs

Companies using AI-assisted project management tools report an average 15 percent improvement in project delivery productivity.¹⁰ On the specific task of monthly reporting — which typically consumes 3–5 working days per reporting period for a major project’s controls team — consistent findings indicate that AI assistance reduces this by 40–60 percent. This frees experienced project controls professionals to perform root cause analysis, contractor engagement, and recovery planning rather than data assembly and chart production.

3.2 Conservative Financial Model

The following model is intentionally conservative and uses a single EUR 1 billion project as the unit of analysis:

Value Driver	Basis	Annual Value (EUR)
Earlier detection — 1 avoided 2% cost growth	One intervention per project year	20,000,000

¹⁰Various sources aggregated in: Celoxis (2025). “AI and Machine Learning in Project Management.” 15% average productivity improvement; 61% on-time delivery with AI tools vs. 47% without.

Value Driver	Basis	Annual Value (EUR)
Reporting efficiency — 2 FTE-months freed per year	Senior engineer at EUR 15k/month fully loaded	360,000
Consistent escalation — avoided one late crisis	50% probability × EUR 5M average cost of late recovery	2,500,000
Total (conservative)		~22,860,000

Against this, the implementation costs are modest:

Cost Item	Basis	Annual Cost (EUR)
Claude Code enterprise licences (10 users)	Current enterprise pricing	~60,000
Internal implementation effort (one-time)	15 person-days × senior engineer rate	~30,000
Ongoing maintenance and data governance	0.25 FTE	~75,000
Total		~165,000 p.a.

The implied return on investment exceeds 100:1 on conservative assumptions, driven almost entirely by the value of a single avoided late-stage intervention. The ratio improves further on a portfolio of projects, where the fixed infrastructure cost (licences, governance, templates) is shared.

3.3 Contextual Benchmark

Companies that use AI-driven tools in project management deliver 61 percent of their projects on time, compared to 47 percent for those that do not — a 14-percentage-point improvement.¹¹ On a programme of five concurrent capital projects, each averaging EUR 500 million in CAPEX, moving one project from the “late” to “on-time” bucket — with a typical delay cost of 5–10 percent of CAPEX — represents EUR 25–50 million in value creation from the portfolio uplift alone. The open-source code base eliminates implementation risk as a financial objection: the system

¹¹Ibid.

can be inspected in detail before deployment, extended without vendor dependency, and discontinued without sunk cost if the pilot does not demonstrate value.

3.4 AI-Specific Risks

Risk: AI produces confident but incorrect analysis

Likelihood: Medium. Consequence: Medium.

The system mitigates this through schema validation, provenance requirements on all heuristics, and mandatory cross-checking of EAC against parametric benchmarks. No AI output is presented as final without a named human reviewer. The system generates recommendations; it does not make decisions.

Risk: Sensitive commercial data processed through third-party AI systems

Likelihood: Low with controls. Consequence: High.

Claude Code can be deployed on-premises or in a private cloud environment using Anthropic's enterprise API. No project data need leave the organisation's controlled infrastructure. The data governance policy establishes which data categories are permissible and in what environment.

Risk: Over-reliance reduces human expertise over time

Likelihood: Low with design. Consequence: Medium.

This risk is real and documented in analogous automation contexts. The mitigation is design: the system explicitly requires human interpretation of every report, ensures that escalation recommendations are reviewed not auto-executed, and generates plain-language explanations that build rather than bypass analytical understanding.

3.5 Organisational Risks

Risk: Resistance from project controls professionals

Likelihood: High without change management. Consequence: Medium.

The most effective mitigation is framing. This system removes the least engaging portions of a cost engineer's work — data assembly, chart production, threshold checking — and returns time for the high-judgment work that experienced professionals value. Early engagement of senior project controls staff in the pilot design, and visible credit for the improved reporting outputs, is the primary change management lever.

Risk: Data quality too poor for meaningful analysis

Likelihood: Medium. Consequence: Medium.

The system will surface data quality problems that were previously obscured by manual report production. This is a feature, not a bug. A data quality review before the system goes live on the pilot project establishes a clean baseline and closes the gaps that would otherwise persist undetected.

Risk: Pilot succeeds but rollout stalls

Likelihood: Medium without board sponsorship. Consequence: High.

The pattern of successful pilot — stalled rollout — is the dominant failure mode for enterprise technology adoption. The mitigation is board-level ownership of the rollout mandate and a defined CAPEX threshold above which the system is mandatory, not optional.

3.6 What This System Does Not Do

It is equally important to state explicitly what this system does not replace:

- It does not replace the Project Director's judgment on whether to escalate a situation to the board
- It does not negotiate with contractors, assess claim merits, or draft legal correspondence without human review
- It does not produce FID-quality cost estimates; it analyses cost *performance* against an already-established baseline
- It does not replace the physical site inspection, quality review, or HSE incident investigation
- It does not guarantee project success; it improves the information available to the humans who determine outcome

Projects succeed because of skilled, experienced, well-led teams. This system makes it harder for those teams to miss early warning signals, easier to produce consistent reporting, and more likely that the right information reaches the right decision-maker at the right time.

Appendix A — Key Metrics Produced by the EVM Module

Metric	Formula	What it tells management
CPI	$BCWP \div ACWP$	Cost efficiency: EUR earned per EUR spent
SPI	$BCWP \div BCWS$	Schedule efficiency: earned value vs plan
EAC	$ACWP +$	Final cost forecast, accounting for schedule pressure
(Composite)	$(BAC - BCWP) \div (CPI \times SPI)$	
VAC	$BAC - EAC$	Expected over- or under-run at completion
TCPI	$(BAC - BCWP) \div (BAC - ACWP)$	CPI required for all remaining work to finish within budget; >1.10 indicates re-baselining is needed

RAG thresholds (industry standard):

Index	GREEN	AMBER	RED
CPI	≥ 0.95	$0.85 - 0.94$	< 0.85
SPI	≥ 0.95	$0.85 - 0.94$	< 0.85

Appendix B — Technical Resource

The working system — EVM skill, scaffolding templates, domain heuristics, JSON schemas, specialist agent definitions, and a complete example project — is available at:

https://github.com/jmeier1963/large_capital_project_management

The repository includes: - `evm/evm_calculator.py` — EVM engine (713 lines, no dependencies beyond pandas and matplotlib) - `budget-estimator/budget_estimator.py` — parametric CAPEX engine with P50/P90 Monte Carlo - `scaffolding/` — project configuration templates for immediate deployment - `examples/H2-PIPE-DE-001/` — complete worked example: 500 km H2 pipeline, five months of EVM data, illustrated analysis outputs - `README.md` — step-by-step installation and usage instructions

Installation requires copying three files and running one `pip install` command.

References
